

Just like before, your homework submission must contain the following to receive full credit:

- A single script called “hmkw5.m” or “hmkw5.py” that requires no arguments and prints the things (and makes the plots) I ask for in bold below. Each output should be labeled.
- A pdf called hmkw5\_results.pdf with your solutions to all theory problems and also all the PLOTS generated in the hmkw5 script. See the posted homework example.
- As many other code files as you wish - but only things you wrote or solutions to previous homeworks.

1. Consider the quadratic image denoising model

$$\text{L2 denoising: } \min \quad \frac{\mu}{2} \|\nabla x\|^2 + \frac{1}{2} \|x - b\|^2$$

where  $\mu$  is a scaling parameter and  $b$  is a noisy image. We use  $\|\cdot\|$  to denote the  $\ell_2$ /Frobenius norm. Note that  $\nabla$  denotes the discrete gradient (i.e., your `grad2d` function). For this problem, you should choose  $b$  to be a noisy, black-and-white test image of your choice, such as Barbara, Lena, the Shepp-Logan phantom, etc... Scale the image to have pixels between 0 and 1, and set  $\mu = 2$ .

- Write the optimality condition for this problem by taking the gradient and setting it equal to zero. Now, put the optimality condition in the form of a large linear system  $Ax = b$ .**
- Create a function or function handle that computes the linear operator  $A$  from part (a). The only argument to this function should be a 2d pixel array representing  $x$ . This can be done with only 1 line of code when using your `grad2d` and `div2d` functions. Both the inputs and the outputs of your function **MUST** be 2d arrays.
- Write a function with signature

```
function x, resids = richardson(A,b,x,t)
```

that solves the system  $Ax = b$  using Richardson iteration with stepsize  $t$ , where  $A$  is *any* arbitrary handle/function that computes a symmetric positive definite linear operator. The method must accept an initial starting point  $x$  of arbitrary dimensions (i.e., 1d, 2d, 3d, etc...). The method should terminate when the residual norm falls below  $10^{-6}$ . The method returns the solution,  $x$ , and also a vector containing the residual norms on every iteration. Use your code to minimize the L2 denoising energy above. **Make a plot of residual norms vs iteration number. Display the noisy image beside the denoised image. Report how many iterations this took.**

- Write a function

```
function x, resids = conjgrad(A,b,x)
```

that implements the conjugate gradient method for solving  $Ax = b$  where  $A$  is a symmetric positive definite linear operator. The method must accept an initial guess  $x$  of arbitrary dimensions, and return the solution to the system in addition to a vector of residuals. Use this routine to solve the L2 denoising problem. **Make a plot of residual norms vs iteration number. Display the noisy image beside the denoised image. Report how many iterations this took.**

- (e) Now, set  $\mu = 10$ , and re-run your experiments. **Show your images and convergence plots. What happened to your iteration counts? Why did this happen? Make a mathematical argument with some equations.**
- (f) Write a function

```
function x = l2denoise(b,mu)
```

that directly computes the exact solution (up to numerical rounding errors) of the denoising problem using Fourier transforms (i.e., without using an iterative method). Your implementation can call `fft2/iff2` at most 4 times total, and you may not use any loops of any kind. **Call your method to solve the denoising problem, and evaluate and report the norm of the gradient of the objective function. This value should be very small.**

2. In this problem you'll do spectral clustering on a toy dataset. Download the “two moons” dataset generator from the website.

- (a) Build a dataset with 100 points. Scatterplot the point and have a look at them. Pick some reasonable value of  $\sigma$  (the length scale for the problem), and build the  $100 \times 100$  pairwise similarity matrix  $S$  with

$$S(i, j) = \exp(-\|x_i - x_j\|/\sigma).$$

Your value of  $\sigma$  should be small enough that the two moons are “far apart” on this distance scale.

- (b) Compute the diagonal normalization matrix with  $D(i, i) = \sum_j S(i, j)$ . This matrix just contains the row sums of the similarity matrix. Use this matrix to compute the normalized similarity matrix

$$\hat{S} = D^{-1/2} S D^{-1/2}.$$

Note that we split the normalization matrix up and apply it to both sides of  $S$  to preserve symmetry.

- (c) Compute the eigenvalue decomposition of  $\hat{S}$ , and select the second eigenvector  $u_2$  (i.e., the eigenvector with second largest eigenvalue). This is the “spectral embedding” of your data points. **Plot the vector  $u_2$ .** The two moons should now look linearly separable (expect for maybe a few points).
- (d) Run k-means clustering on the entries in  $u_2$  (I suggest you use the `kmeans` command in Matlab, and `sklearn.cluster.KMeans` in Python) to produce 2 clusters. **Plot the original datapoints with two different colors to indicate which class the data-point was assigned by k-means. You should have separated the moons almost perfectly.**

3. Repeat Question 2, but this time with 100,000 datapoints. The process outlined below is described in detail in the paper “Spectral Grouping Using the Nystrom Method.”

- (a) Build a Nystrom approximation of  $S$  of the form  $S_{ny} = CW^{-1}C^T \approx S$  by sampling 200 randomly chosen columns of  $S$ .
- (b) We can't compute row sums of  $S$  because it's too big to form explicitly, so instead compute the row-sums of  $S_{ny}$ , and define the diagonal matrix  $D(ii) = \sum_j S_{ny}(i, j)$ . Normalize the rows/columns you sampled to find the matrix  $C_n$  containing 200 column samples from the normalized matrix  $\hat{S}_{ny} = D^{-1/2}SD^{-1/2}$ . Define  $W_n$  by sampling the corresponding rows from  $C_n$ . Define the matrix  $M_n$  comprising all the rows of  $C_n$  corresponding to un-sampled column indices.
- (c) Compute the approximate eigenvectors of  $\hat{S}_{ny}$ . Do this in stages:

Compute orthogonalization matrix:  $\hat{W} = W_n + W_n^{-1/2}M_n^T M_n W_n^{-1/2}$

Compute eigen-decomposition:  $U_W D_W U_W^T = \hat{W}$

Compute approximate eigenvectors:  $U = C_n W_n^{-1/2} U_W D_W^{-1/2}$

Rescale:  $U(:, k) \leftarrow U(:, k) / U(:, 200)$ , for  $k = 1, 2, \dots, 200$ .

The factorization  $\hat{S}_{ny} = U D_W U^T$  is now an eigenvalue decomposition. Finally, I recommend that you perform this rescaling step:

Rescale:  $U(:, k) \leftarrow U(:, k) / U(:, 200)$ , for  $k = 1, 2, \dots, 200$ .

Note, the last step of the process (“Rescale”) divides every non-leading eigenvector by the leading eigenvector, i.e., I assume your vectors are ordered so that  $U(:, 200)$  is the vector of largest eigenvalue.

Let  $U_2$  contain the 2nd most significant eigenvector. **Plot**  $U_2$ . The two moons should now be linearly separable.

- (d) Do k-means on  $U_2$ . **Plot the original datapoints with two different colors to indicate which class the datapoint was assigned by k-means. You should have separated the moons almost perfectly.**